LEIBNIZ UNIVERSITÄT HANNOVER

FAKULTÄT FÜR ELEKTROTECHNIK UND INFORMATIK
INSTITUT FÜR VERTEILTE SYSTEME

Multi-purpose Library of Recommender System Algorithms for the Item Prediction Task

Bachelor Thesis

eingereicht von

JULIUS KOLBE

am 11. Juni 2013

Erstprüfer	:	Prof. Dr. techn. Wolfgang Nejdl
Zweitprüfer	:	Jun.-Prof. Dr. rer. nat. Robert Jäschke
Betreuer	:	Ernesto Diaz-Aviles, M. Sc.

EHRENWÖRTLICHE ERKLÄRUNG

Hiermit versichere ich, die vorliegende Bachelor Thesis ohne Hilfe Dritter und nur mit den angegebenen Quellen und Hilfsmitteln angefertigt zu haben. Alle Stellen, die wörtlich oder inhaltlich aus den Quellen entnommen wurden, sind als solche kenntlich gemacht worden. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

Hannover, den 11. Juni 2013

Julius Kolbe

ABSTRACT

In the context of this thesis, we implemented a recommender system library called *recsyslab* [13]. We used the programming language Python [12] to this end, given its simplicity and elegance. The goal of recsyslab is to provide easy access to current recommendation algorithms, by being easy to use and having easily readable source code. The development of recsyslab is the major contribution of this thesis. This document is intended to document its design and facilitate its use.

ZUSAMMENFASSUNG

Im Rahmen dieser Bachelorarbeit haben wir eine Bibliothek namens *recsyslab* [13] in der Programmiersprache Python [12] implementiert. Wir haben Python aufgrund dessen Einfachheit und Eleganz benutzt. Das Ziel von recsyslab ist es, einen möglichst einfachen Zugang zu Empfehlungsalgorithmen zu gewähren, indem es einfach zu benutzen ist und der Quellcode gut zu lesen ist. Die Entwicklung von recsyslab ist der Hauptteil der Bachelorarbeit. Dieses Dokument soll ihre Struktur dokumentieren und ihren Gebrauch vereinfachen.

CONTENTS

INTRODUCTION

Recommender systems are widely used in e-commerce and are an active field of research; new and refined algorithms are proposed continuously. Easy and comprehensive evaluation tools are required to assess their performance in a systematic way.

With the simple syntax and the interactivity of Python, *recsyslab*, which is available at

> https://github.com/Foolius/recsyslab/archive/master.zip,

is designed in such a way, so as to facilitate experimentation with different recommender systems algorithms. Our goal is to provide beginners and experts with an extensible tool to evaluate classic and modern algorithms, to comprehend them, and to extend the existing core implementation if necessary. With this document, recommender systems' enthusiasts should be able to get started in trying out different recommender algorithms without needing any previous knowledge or long configuration.

1.1 WHAT A RECOMMENDER SYSTEM DOES

A Recommender System works in a scenario with users, items, and interactions between these two. Such a scenario could be an online shop such as amazon.com, where the interactions are purchases of items by users or a video platform, where the users interact with items (videos) by watching them, e.g., youtube.com. Based on the past interactions of the users a recommender system presents them with a ranked list of items, which meets their interest and taste.

Two fundamental tasks for recommender systems are:

- **Rating Prediction**: When the feedback is provided explicitly through ratings (e.g., stars in a 1 to 5 scale), the scenario is called rating prediction, and the task consists of predicting the real value of the rating. This task can be cast as a regression problem.

  The rating prediction task gained popularity when Netflix announced a prize [10] in September 2009, offering a one-million-dollar prize to a team able to improve Netflix's existing recommender system.

- **Item Prediction**: When the interactions are implicit, like, in case of purchases or clicks, then the problem can be cast as item prediction or top-N recommendation. In this case, the task is to recommend a short list of items to meet the particular taste of the users. Please note that in this scenario we are not interested in regressing a value for a rating, but in relative ranking of items that meet the user's individual information needs.

In this work, the focus lies precisely on this scenario of implicit feedback and item prediction. The algorithms implemented within recsyslab compute a score which is used to rank the items relative to each other, but this score is not an absolute rating.

Note that ratings can also be interpreted as the strength for implicit feedback. Some algorithms implemented in this library make use of this information, but none of them explicitly predict ratings, unlike the usual practice in rating prediction scenarios. Rating information is used to derive a relative order within the items, that is, it is sufficient to only have a ranked list of items according to the preferences of the users. This information is then used to train our item prediction algorithms.

1.2 CONTRIBUTIONS

The contributions of this bachelor thesis are:

1. The implementation of an easy to use and Open Source library: *recsyslab*, which is released under the GNU GPLv3[1] and includes a set of the most important state-of-the-art recommender system algorithms for item prediction. Special care has been given to the source code produced, to keep it simple and readable, and as close as possible to the pseudocode of the algorithms.

2. Supporting infrastructure for testing recommender algorithms.

3. Extensive user manual.

4. Explanations of the technical background and software design.

1.3 STRUCTURE OF THIS DOCUMENT

The next chapter includes an overview of the research area in recommender systems, together with explanations of different evaluation metrics and the leave-one-out protocol commonly used to evaluate top-N recommender algorithms.

After that, we provide short descriptions of already existing projects which provide something similar to recsyslab (Chapter 3).

In Chapter 4, we present the implemented recommender algorithms along with the core of their implementation.

After that, we show the inner structure of recsyslab to make it easy for developers to extend or change the library (Chapter 5).

Chapter 6 provides an extensive guide for recsyslab explaining how the library is used.

The last chapter includes conclusions and an outlook to future work (Chapter 7).

BACKGROUND

In this chapter, we provide a short overview of the recommender systems realm, and one particular protocol to evaluate them.

2.1 COLLABORATIVE FILTERING

Broadly speaking, there are two approaches for recommender systems:

- Collaborative filtering

- Content-based filtering

Collaborative filtering is, in general, the process of finding information or patterns in large datasets to predict user preferences. By assuming that users who have the same opinion on some items will also agree on others, collaborative filtering models are capable of generating predictions without using additional information like the age of the users [3].

Content-based filtering, on the other hand, uses information about users and items. For example, the age of the users or the name or genre of the items [4].

This work focuses only on collaborative filtering. So when we are speaking about recommender systems we mean recommender systems using the collaborative filtering approach.

2.2 MATRIX FACTORIZATION

An important class of recommender algorithms for collaborative filtering uses matrix factorization. Matrix factorization approximates a large and sparse matrix into a product of two low-rank matrices [8].

In recommender systems, the matrix representing the dataset of user-item interactions (M) is factorized into two matrices W and H such that $\hat{M} = W\,H^{\top}$. The matrices W and H are supposed to represent abstract features of each item and user correspondingly. For recommendation, the dot product of the feature vector of an user and an item gives a score, with which we can sort the items and recommend the most suitable ones.

2.3 EVALUATION METHODS

To evaluate a recommender algorithm we have to split up the database into one for training and one for evaluation. There are different methods to split the database, but in the library only one is implemented, which is the leave-one-out protocol [6]. One can use the leave-one-out protocol with many different metrics which are also explained here.

2.3.1 *Leave-one-out Protocol*

The leave-one-out protocol works as follows:

1. Randomly choose one interaction for every user. These are the hidden interactions.

2. Put the hidden interactions into the test set.

3. Put all other interactions into the training set.

4. Let the recommender recommend the first N items for every user by ranking all items, except the ones the user already interacted with.

5. If the hidden interaction of the user is in the recommendations for this user, the recommender gets a hit.

6. Count the hits and record the position of the interaction in the recommendations.

7. Use this information to compute the various metrics.

Because of the split, we can test the recommender on every user. This would not be possible if we would, for example, split by randomly choosing 1% of the interactions. Then it would be likely that there are some users without a hidden item. And without a hidden item it would be impossible to measure the quality of the recommendations for this user.

2.3.2 *Evaluation metrics*

These are a selection of different metrics to rate the recommendations. By default the evaluations are executed with only one hidden item. But generally, the metrics should also work with more than just one.

For the notation: U is the set of users, H is the set of hidden items and H_u is the set of hidden items for user u. T is the set used for training. $TopN_u$ is the set of top N recommendations for user u so the number of items the recommender is allowed to recommend is N. For the metrics, where the order in which the items are recommended count, $TopN_u$ is a list sorted by score in decreasing order. To get an implicit score of each item the recommender recommends all items.

The recommender algorithm recommends N items. This N has to be chosen according to the deployment scenario. For example, the number of items that can get recommended at once depends on the layout of the website. According to a chosen N and with the metrics of recsyslab, the right algorithm can be found.

Except for AUC and MRHR, the metrics, defined later in this section, do not take the position of the items in the recommendation list into account.

Whether the position is important or not, depends on the use case. When all recommended items are presented equally well to the user, the position of the items in the recommendation is not as important as whether or not

the right items are recommended. But in contrast to this, when the user is shown a long list of items and it is likely that the user will not scroll down until the end of the list, one would rather care about the position of the items in the list of recommendations.

2.3.2.1   *Hitrate/Recall@N*

This metric lets the recommender recommend N items. If the hidden item is under the N recommended items, the recommender gets a hit [21, 32]. So the Recall@N is the fraction of users who get recommended a relevant item when the recommender can recommend N items. So the hitrate is

$$\text{Recall@N} = \frac{\sum_{u \in U} H_u \cap \text{topN}_u}{|H|}. \tag{2.1}$$

This metric is very intuitive. For example, imagine that you show the user 10 items, then Recall@10 would be the chance of showing the user an item he will interact with. But this metric does not take the number of recommended items into account.

2.3.2.2   *Precision*

The precision [32] is defined as

$$\text{Precision} = \frac{\sum_{u \in U} H_u \cap \text{topN}_u}{N \times |U|}. \tag{2.2}$$

As one can clearly see, this metric takes the number of recommended items into account. This will probably lead to worse results as the number of recommended items increases.

2.3.2.3   *F1*

The F1 metric [32] tries to balance hitrate and precision by taking both into account. It is defined as

$$\text{F1} = \frac{2 \times \text{Recall@N} \times \text{Precision}}{\text{Recall@N} + \text{Precision}}. \tag{2.3}$$

2.3.2.4   *Mean Reciprocal Hitrate*

The mean reciprocal hitrate (MRHR), or the more general mean reciprocal rank [27], counts the hits but punishes them more, the lower they appear in the list of recommendations. So if the hidden item appears first in the list of recommendations, the hit counts as one. But when it is in the second, position the hit counts only as one half and so on. It is defined by

$$\text{MRHR} = \frac{1}{|U|} \sum_{u \in U} \frac{1}{\text{pos}(\text{topN}_u, H_u)}, \tag{2.4}$$

where N is the number of items in the dataset and $pos(topN_u, H_u)$ is the position of the hidden item in the recommendation.

2.3.2.5 *Area under the ROC (AUC)*

AUC [29] counts the number of items that the recommender rates higher than the hidden item, normalized by the number of items the recommender can rate higher. This is in turn summed up for every user and normalized by the number of users.

To get an implicit score of each item, the recommender recommends all items in a list, sorted by score in a decreasing order. This is in fact the same as for the other metrics, except that the recommender can recommend as many items as possible. It is defined as

$$AUC = \frac{1}{|U|} \sum_{u \in U} \frac{1}{|E(u)|} \sum_{(i,j) \in E(u)} \delta(x_{ui} > x_{uj}), \tag{2.5}$$

where x_{ui} is the predicted score of the interaction between User u and item i. δ is defined as follows

$$\delta(x) = \begin{cases} 1, & \text{if x is true,} \\ 0, & \text{otherwise.} \end{cases} \tag{2.6}$$

And $E(u)$ is

$$E(u) = \{(i,j) | (u,i) \in H \wedge (u,j) \notin (H \cup T)\}. \tag{2.7}$$

RELATED WORK

There is a wide range of projects providing implementations for recommender system. Some of them are described in this chapter to give a quick overview and comparison.

3.1 MYMEDIALITE

MyMediaLite [19] is an open source project developed at the University of Hildesheim and provides several algorithms for rating prediction and item prediction. It is written in C# and can be used with a command line interface. It also provides a graphical interface to demonstrate recommender algorithms.

3.2 PREA (PERSONALIZED RECOMMENDATION ALGORITHMS TOOLKIT)

PREA [23] is an open source project written in Java. It provides a wide range of recommender algorithms and evaluation metrics to test them. It is maintained by the Georgia Institute of Technology.

3.3 APACHE MAHOUT

Mahout [7] is an open source library in Java. It is implemented on top of Apache Hadoop, so it uses the map/reduce paradigm. This means it can run on different, independent computers.

3.4 DUINE FRAMEWORK

The Duine Framework [5] is an open source project written in Java by the Telematica Instituut/Novay. The recommender of the Duine Framework combines multiple prediction techniques to exploit the strengths of the different techniques and to avoid their weaknesses.

3.5 COFI

Cofi [2] provides an algorithm for the rating prediction task called Maximum Margin Matrix Factorization. It is open source and written in C++.

3.6 LENSKIT

Lenskit [18] is a toolkit which provides several recommender algorithms and an infrastructure to evaluate them. It is an open source project by the University of Minnesota.

3.7 COMPARISON

The table below compares the algorithms implemented by the different frameworks.

Category	Feature	PREA	Mahout	Duine	Cofi	MyMedia	**recsyslab**
Baselines	Constant	O			O	O	O
	User/Item Average	O	O	O	O	O	
	Random	O				O	O
Memory-based CF	User-based CF [34]	O	O	O	O	O	O
	Item-based CF [26]	O	O	O	O	O	O
	Default Vote, Inf-User-Frex [16]	O	O				
	Slope-One [24]	O	O			O	O
Matrix Factorization	SVD [28]	O	O		O	O	O
	NMF [33]	O					
	PMF [31]	O					
	Bayesian PMF [30]	O					O
	Non-linear PMF [22]	O					
	RankMFX [17]						O
Other methods	Fast NPCA [36]	O					
	Rank-based CF [35]	O					
Evaluation Metric	(N)MAE	O	O	O	O	O	
	RMSE	O	O		O	O	
	HLU/NDCG	O				O	
	Kendall's Tau, Spearman	O					
	Precision/Recall/F1		O			O	O
	ARHR/MRHR						O
Miscelaneous	Sparse Vector/Matrix	O	O	O	O	O	O
	Item Recommender for Positive-only Data					O	O
	Release Year	2011	2005	2009	2004	2009	2013
	Language	Java	Java	Java	C++	C#	Python
	License	GPL	LGPL	LGPL	GPL	GPL	GPLv3

Table 3.1: Comparison of different recommender frameworks

RECOMMENDATION ALGORITHMS

In this chapter we provide an overview of how the algorithms we have implemented work. To illustrate the simplicity of the code, we also show the core part of the source code of the respective recommender algorithm in recsyslab. For further explanations please refer to the cited papers.

4.1 NON-PERSONALIZED ALGORITHMS

In this section we describe two very simple and basic recommendation algorithms, that we have implemented for comparison with the more sophisticated algorithms.

4.1.1 *Constant*

The constant recommender algorithm counts the number of interactions for each item and sorts this in a decreasing order of interactions. Then it recommends the top items of this list. So it recommends the items which are the most popular over all users and does not account for any personalizations. The intuition here is that the most popular items will be interesting for everyone. However, the results show that algorithms, which personalize the recommendation based on the interaction history of the users, perform much better. The pseudocode of this algorithm would look like this:

```
countDict = dict initialized with all ItemIDs as keys and 0 as values

for interaction in dataset:
    countDict(ItemID(interaction))++

ranking = list(ItemIDs in decreasing order of their values)

return first N items of ranking
```

In python it looks like this:

```python
def __init__(self, dbdict):
    self.dictionary = {}
    self.sortedList = []
    for data in dbdict.iteritems():
        for item, rating in iter(data[1]):
            if item in self.dictionary:
                self.dictionary[item] += rating
            else:
                self.dictionary[item] = rating

    self.sortedList = helper.sortList(self.dictionary.iteritems())
```

```python
def getRec(self, user, n):
    return self.sortedList[:n]
```

4.1.2 *Random*

The random recommender algorithm chooses items to recommend randomly. Even though this algorithm recommends different items for different users, it is not personalized, because the recommendations are independent of the previous interactions of the user. The pseudocode of this algorithm would be:

```
return N randomly chosen items
```

And in Python:

```python
def __init__(self, dbdict, seed):
    self.maxIid = 0
    self.seed = seed
    for data in dbdict.iteritems():
        for itemRating in iter(data[1]):
            item = itemRating[0]
            if item > self.maxIid:
                self.maxIid = item
    self.maxIid += 1

def getRec(self, user, n):
    random.seed(self.seed)
    if self.maxIid < n or n == -1:
        l = range(self.maxIid)
        random.shuffle(l)
        return l
    return list(random.sample(range(self.maxIid), n))
```

4.2 K-NEAREST-NEIGHBOR

This class of recommendation algorithms works by searching neighbors of either items or users based on a similarity function, which is the cosine in this library. The similarity function is interchangeable. For example in [21], two similarity functions are compared. The cosine similarity performs best, so in recsyslab only the cosine similarity is implemented. For two vectors $\vec{v}, \vec{u}$ the cosine similarity is defined as follows:

$$\cos(\vec{v}, \vec{u}) = \frac{\vec{v} \cdot \vec{u}}{||\vec{v}||_2 ||\vec{u}||_2} \tag{4.1}$$

With the similarity it is possible to rank items. In the item-based algorithm, for each item, we compute the sum of similarities with the items the user already bought and rank accordingly. In the user-based algorithm, for every item, we sum up the similarities of the users who bought this item and rank according to this score.

4.2.1 *Item-Based*

For this algorithm, the database has to be represented as a matrix where the rows correspond to the users and the columns to the items. Then the entry (i,j) represents the number of transactions that occured between the i-th user and the j-th item.

The algorithm interprets the columns of the matrix, i.e. the items as vectors and computes the similarities by computing their cosine. To build the model, for each item, the algorithm computes the n most similar items. The pseudocode:

```
for every item i
    for every item j
        sim[i,j] = similarity between i and j

for every item i
    for every item j
        if sim[i,j] not one of the n largest in sim[i]
            sim[i,j] = 0
```

In Python it looks like this:

```python
def __init__(self, userItemMatrix, n):
    self.itemUserMatrix = userItemMatrix.transpose()
    self.sim = computeCosSim(self.sim, self.itemUserMatrix)
    self.userItemMatrix = userItemMatrix

    order = self.sim.argsort(1)

    # for each row in sim:
    # Set all entries to 0 except the n highest
    for j in xrange(0, self.sim.shape[1]):
        for i in xrange(0, self.sim.shape[1] - n):
            self.sim[j, order[j, i]] = 0
```

`computeCosSim` looks like this:

```python
def computeCosSim(sim, matrix):
    for i in xrange(1, sim.shape[1]):

        for j in xrange(0, i):
            sim[i, j] = sim[j, i] = cos(
                matrix[i], matrix[j])
    return sim
```

This function is also used in the user based k-Nearest-Neighbor.

To compute recommendations for user U, the algorithm computes the union of the n most similar items for each item that U interacted with. From this set, the items that U already interacted with are removed. For each item remaining in this set, we compute the sum of its similarities to the items U interacted with. Finally, these items are sorted in a decreasing order of the sum of similarities, and the first n items will be recommended [21]. The pseudocode is:

```
for every item i user u bought
    itemSimVector = itemSimVector + vector of i

for every item i user u bought
    itemSimVector[i] = 0

return the N items with the highest value in itemSimVector
```

The Python equivalent:

```python
def getRec(self, u, n):
    # x are the similarities of each item to the items u bought
    # w.r.t. only the highest n similarities are saved.
    x = self.userItemMatrix[u] * self.sim

    # Throw out items the user already purchased
    for i in xrange(0, self.sim.shape[0]):
        if self.userItemMatrix[u, i] != 0:
            x[0, i] = 0

    order = x.argsort()
    l = []
    for i in xrange(1, n + 1):
        l.append(order[0, -i])
    return l
```

4.2.2   *User-Based*

The user-based k-Nearest-Neighbor [14] is very similar to the item-based algorithm. But instead of interpreting the columns as vectors, we interpret the lines or users of the matrix as vectors, and compute their similarities to other users. Pseudocode:

```
for every user u
    for every user o
        sim[u,o] = similarity between u and o

for every user u
    for every user o
        if sim[u,o] not one of the n largest in sim[u]]
            sim[u,o] = 0
```

Python:

```python
def __init__(self, userItemMatrix, n):
    self.userItemMatrix = userItemMatrix
    self.sim = np.zeros((userItemMatrix.shape[0], userItemMatrix.shape
        [0]))
    self.sim = computeCosSim(self.sim, self.userItemMatrix)

    order = self.sim.argsort(1)

    # for each row in sim:
```

```
    # Set all entries to 0 except the n highest
    for j in xrange(0, self.sim.shape[1]):
        for i in xrange(0, self.sim.shape[1] - n):
            self.sim[j, order[j, i]] = 0
```

Then for each item i, we sum up the similarities between U and the users who interacted with i. Again, we remove all the items that U already interacted with, sort in a decreasing order of the sum and recommend the first n items. Pseudocode:

```
for every item i
    for every user o
        score(i) = score(i) + sim[U,o]

for every item i user u bought
    score(i) = 0

return the N items with the highest value in score
```

Python:

```python
def getRec(self, u, n):
    # x is the weighted sum of the items
    # weighted with the similarity between u and the
    # other users.
    x = self.sim[u] * self.userItemMatrix

    for i in xrange(0, self.sim.shape[0]):
        if self.userItemMatrix[u, i] != 0:
            x[0, i] = 0

    order = x.argsort()
    l = []
    for i in xrange(1, n + 1):
        l.append(order[0, -i])

    return l
```

4.3 MATRIX FACTORIZATION

As mentioned in Section 2.2, the matrix factorization techniques generate two matrices W and H, such that $\hat{M} = W\,H^{\top}$. Each of the implemented algorithms train these two matrices with stochastic gradient descent. In each iteration the model is trained with:

- a randomly chosen user, U

- a randomly chosen item, I, the user U interacted with, called the positive item, and

- a randomly chosen item, J, that the user U did not already interact with, called the negative item.

The features of U, I and J are trained using the derivative of a loss function.

BPRMF and RankMFX share the code in the module `recommender.mf`. The only difference is the loss function. In pseudocode the model update happening in each iteration looks as follows:

```
U = randomly chosen user
I = randomly chosen item U interacted with
J = randomly chosen item U did not interact with

X=H[i] - H[j]
wx = dot product of W[u] and X
dloss = (derivative of the loss function of wx and 1) * learningRate

W[u] += dloss * (H[i] - H[j]) # These three lines
H[i] += dloss * W[u]          # have to be
H[j] += dloss * -W[u]         # executed at once
```

The Python code doing the same looks like this:

```python
u = random.choice(R.keys())

userItems = [x[0] for x in R[u]]
# the positive example
i = userItems[np.random.random_integers(0, len(userItems) - 1)]
# the negative example
j = np.random.random_integers(0, m_items)
# if  j is also relevant for u we continue
# we need to see a negative example to contrast the positive one
while j in userItems:
    j = np.random.random_integers(0, m_items)

X = H[i] - H[j]
wx = np.dot(W[u], X)
dloss = dlossF(wx, y)

# temp
wu = W[u]
hi = H[i]
hj = H[j]

if dloss != 0.0:
    # Updates
    eta_dloss = learningRate * dloss
    W[u] += eta_dloss * (hi - hj)
    H[i] += eta_dloss * wu
    H[j] += eta_dloss * (-wu)

    W[u] *= scaling_factorU
    H[i] *= scaling_factorI
    H[j] *= scaling_factorJ
```

For the recommendations, we rank the items with a score. The score of an item for a user is the dot product of the feature vector of the user and the

feature vector of the item. All matrix factorization algorithms use this. In pseudocode:

```
for every item i
    scoreList(i) = W[u] * H[i]

sortByScore(score)
return first N of score
```

And in Python:

```python
def getRec(self, u, n):
    scoredict = {}
    for i in range(0, self.H.shape[0]):
        if not i in [x[0] for x in self.R[u]]:
            scoredict[i] = np.dot(self.W[u], self.H[i])

    if n == -1:
        n = len(scoredict)
    import util.helper
    return util.helper.sortList(scoredict.iteritems())[:n]
```

4.3.1 *BPRMF*

In our implementation of BPMRF (Bayesian Personalized Ranking with Matrix Factorization) we use the logLoss to train the model. The logLoss is defined as

$$\text{logLoss}(a, y) = \log(1 + \exp(-ay)). \tag{4.2}$$

And the derivative of the logLoss is

$$\frac{\partial}{\partial y}\text{logLoss}(a, y) = -\frac{a}{\exp(ay) + 1}. \tag{4.3}$$

For further information please refer to [29].

4.3.2 *RankMFX*

RankMFX uses the hingeLoss. It is defined as

$$\text{hingeLoss}(a, y) = \max(0, 1 - ay). \tag{4.4}$$

And its derivative is

$$\frac{\partial}{\partial y}\text{hingeLoss}(a, y) = \begin{cases} -y, & ay < 1, \\ 0, & \text{otherwise.} \end{cases} \tag{4.5}$$

See also [17].

4.3.3   *Ranking SVD (Sparse SVD)*

Ranking SVD (Singular Value Decomposition) uses the quadratic loss and the difference between the predicted score of the positive item and the negative, minus the actual score of the positive item [20]. Ranking SVD uses code different from BPRMF and RankMFX, but the iterations work the same way. It also randomly chooses a user and a positive and a negative item in each iteration of the stochastic gradient descent. But the model updates are different. In pseudocode:

```
U = randomly chosen user
I = randomly chosen item U interacted with
J = randomly chosen item U did not interact with

ruipred = W[U] * H[I] # Prediction for item I
rujpred = W[U] * H[J] # Prediction for item J
rui = actual value of the interaction between U and I
ruj = 0

dloss = (ruipred - rujpred) - (rui - ruj)
W[U] = W[U] - learningRate * dloss * (H[I] - H[J])
H[I] = H[I] - learningRate * dloss * W[U]
H[J] = H[J] - learningRate * dloss * -W[U]
```

And in Python this part looks like this:

```python
# Choose an user randomly
u = random.choice(R.keys())
# Choose an item the user interacted with
userItems = [x[0] for x in R[u]]
i = random.choice(userItems)
# Choose an item, the user didn't interacted with
i0 = np.random.random_integers(0, m_items - 1)
while i0 in userItems:
    i0 = np.random.random_integers(0, m_items - 1)

# Prediction for the first item
ruipred = userFeatures[u].dot(itemFeatures[i])
# Prediction for the second item
rui0pred = userFeatures[u].dot(itemFeatures[i0])
rui0 = 0
for r in R[u]:
    if r[0] == i:
        rui = r[1]

dloss = (ruipred - rui0pred) - (rui - rui0)
c = userFeatures[u]
userFeatures -= learningRate * dloss * (
    itemFeatures[i] - itemFeatures[i0])
itemFeatures[i] -= learningRate * dloss * c
itemFeatures[i0] -= learningRate * dloss * (-c)
```

4.4 OTHER

In this section we will describe one last algorithm which does not fit into the categories above.

4.4.1 *Slope One*

The slope one recommendation algorithm computes the differences of interaction intensities between items, and uses these differences to indirectly predict the interaction intensity between users and items without any interaction [25]. The model building in pseudocode:

```
for all interaction pairs i, j where i != j
    diffs[i][j] = i.rating - j.rating
    count[i][j]++
```

And in Python:

```python
def __init__(self, R):
    self.diffs = {}
    self.R = R

    for u in R.keys():
        for i, r in R[u]:
            if not i in self.diffs:
                self.diffs[i] = {}

            for i1, r1 in R[u]:
                if i == i1:
                    continue
                if not i1 in self.diffs[i]:
                    self.diffs[i][i1] = [0.0, 0.0]

                self.diffs[i][i1][0] += r - r1
                self.diffs[i][i1][1] += 1
```

For the recommendations, for every item, we sum up its differences with the items the user rated, together with the rating the user gave these other items. Additionaly, this is multiplied by the number of times the difference occured. This is then normalized by the sum of the occurrences of the differences. The following pseudocode elucidates this.

```
for every item i
    ratingSum = 0
    count = 0
    for every item j user u interacted with
        ratingSum = ratingSum + ((rating u gave j)
                                + diffs[i][j])
                                * count[i][j]
        count = count + count[i][j]

    ratingSum /= count
    listToRecommend.add( (i, ratingSum))
```

```
sortAfterRatingSum(listToRecommend)
return first N of listToRecommend
```

And in Python:

```python
def getRec(self, u, n):
    maxIid = max(self.diffs.keys())
    userItems = [x[0] for x in self.R[u]]
    predictionList = []

    for i in xrange(0, maxIid + 1):
        if i in userItems:
            continue

        ratingSum = 0
        count = 0
        for i1, r in self.R[u]:
            if i in self.diffs:
                if i1 in self.diffs[i]:
                    ratingSum += (r + self.diffs[i][i1][0]
                                  ) * self.diffs[i][i1][1]
                    count += self.diffs[i][i1][1]

        if count == 0:
            continue
        ratingSum /= count
        predictionList.append((i, ratingSum))

    import util.helper
    sortedList = util.helper.sortList(predictionList)

    return sortedList[:n]
```

DESIGN AND IMPLEMENTATION

In this chapter we present a brief overview of the library by explaining its overall structure and how to extend it.

Dependencies

To run recsyslab, one needs a Python 2.7.x interpreter [12] and NumPy [11], a package for scientific computing with Python, to be installed.

Input Dataset Format

At the moment, only one type of dataset is supported. The dataset has to be a textfile where each line is of this form:

```
UserID<separator>ItemID<separator>NumberOfInteractions
```

<separator> is an arbitrary string, but it has to be the same throughout the whole dataset. NumberOfInteractions is optional and can be omitted. In this case one will be assumed. Everything that follows NumberOfInteractions<separator> will be ignored. Please note that when you are omitting NumberOfInteractions but have something else after the ItemId, this will be recognized as NumberOfInteractions.

We recommend to use the MovieLens database (see 6.1) with 100,000 ratings. It is easy to obtain, does not need any modifications to work with our library and has a reasonable size. Additionaly, we will use this dataset in the examples of the user manual in Chapter 6.

5.1 GENERAL STRUCTURE

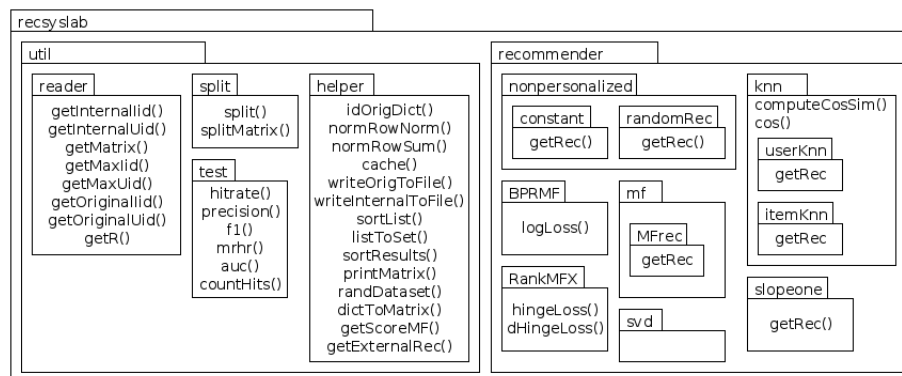Here is a figure illustrating the structure of recsyslab:

Figure 5.1: Structure of recsyslab.

The Util Package

The util package contains several modules for the things happening around the recommender algorithms and the methods for the test metrics. The modules in the util package are:

READER manages the data

SPLIT splits up the dataset following the leave-one-out protocol, see also Chapter 2

HELPER contains several helper functions

TEST contains the methods for computing the test metrics

The `reader` module is a very central part in recsyslab. The class `reader` in the module reader takes care of reading the dataset and supplies the other parts of recsyslab with the kinds of data they need. While the reader reads the dataset, it maps the UserIDs and ItemIDs from the original IDs in the dataset to internal IDs starting at zero and counting up. This guarantees that the IDs are consecutive. When the IDs are consecutive, it is possible to use them as indices in a matrix. Furthermore, the reader constructs a dict and a matrix out of the data using the internal IDs. The dict has the internal UserIDs as keys and the values are tuples representing the interactions of the corresponding user. The first value of these tuples is the internal ItemID of the interaction, the second is the intensity or quantity of the interaction. You can get the dict by calling `getR()` of the `reader` object. The matrix offered by a reader has as many lines as the dataset has users and as many columns as the items that are present in the dataset. The entry at the i-th line and j-th column is then the intensity of the interaction between user i and item j. A zero means that there has not been any interaction

yet. The k-Nearest-Neighbor algorithms need a matrix, like it is provided by a reader. All other recommendation algorithms need a dict of the form described above.

The Recommender Package

The recommender package contains all recommender algorithms. They are separated into the following modules:

NONPERSONALIZED simple, nonpersonalized algorithms

KNN k-Nearest-Neighbors algorithms

MF Pairwise matrix factorization algorithms

BPRMF Bayesian Personalized Ranking Matrix Factorization algorithm

RANKMFX RankMFX algorithm

SLOPEONE Slope One algorithm

SVD Ranking SVD (Sparse SVD)

These modules do not have a `learnModel()` method yet, because the model is trained when the object is instantiated. For descriptions of the algorithms please refer to Chapter 4.

To implement new recommender algorithms which integrate nicely into the infrastructure, the only thing one has to take care of, is to provide a `getRec` function which has an internal UserID as first the first argument and the number of items to recommend as the second argument. Additionaly, it should return a list of internal ItemIDs of the length which has been specified in the second argument.

USER MANUAL

In this chapter we provide a step by step user manual for recsyslab. You can get recsyslab at `https://github.com/Foolius/recsyslab/archive/master.zip`. Simply download the file and unpack it. To run recsyslab you need to install Python [12] and NumPy [11]. Please refer to their websites for information on how to install them.

First, we will explain how to load a dataset and where to get one. Second, we will explain how to use the different recommendation algorithms and how to test them with the provided evaluation metrics. For information about the recommendation algorithms, please refer to Chapter 4. For information about the test metrics refer to Chapter 2.

6.1 DATASETS FOR TESTING

MovieLens [9] is a database provided online by GroupLens, a research lab at the University of Minnesota. One of their research areas is recommender systems and they have built an application where users rate movies and then get recommendations for movies they could potentially like. The MovieLens dataset is a collection of the ratings gathered by this application. For this work, we will interpret the rating as intensity of interaction between users and items. For example, the number of times the user watched a movie.

The dataset is available in three different sizes:

- 100,000 interactions

- 1 million interactions

- 10 million interactions

For the experiments, the smallest dataset is totally sufficient, with the larger datasets the computation time increases largely.

There are also other datasets available, for example the million song dataset [15]. But in the examples below, we use the MovieLens dataset.

6.2 LOAD THE DATASET

The following Python code is also in the file `manual.py`, in the `bin` directory of the library, so you do not have to copy the code out of this PDF document. But if you are working through these examples, please note that you might miss the import of a module when you skip an example, because we do not specify every import needed in each new example. For help, you can also use the inline documentation, which is available as docstrings. You can

display them with the command line utility[1] pydoc. When you are in the bin directory of the library you might call

```
$ pydoc __init__
$ pydoc util
$ pydoc recommender.BPRMF
```

Each of these commands will show the documentation of the specified module. The docstrings are also available in the appendix of this thesis.

In Chapter 5, we explained how the dataset has to look like. For trying things out the MovieLens dataset [9] would work just fine. When you have a dataset, place it into the bin directory of recsyslab and start the python interpreter from the command line with

```
$ python
```

Now import the `util.reader` module and initialize a new reader object with [2]

```
>>> import util.reader
>>> r=util.reader.stringSepReader("u.data","\t")
Start reading the database.
10000 Lines read.
20000 Lines read.
30000 Lines read.
40000 Lines read.
50000 Lines read.
60000 Lines read.
70000 Lines read.
80000 Lines read.
90000 Lines read.
100000 Lines read.
```

Note that it outputs the progress it has already made. The first argument of the constructor is the name of the file containing the dataset, the second the string which is separating the values, here it is a tab. When you are using another dataset, you probably have to change the filename and perhaps also the separating string.

Mapping

The constructor creates a mapping from the original IDs to internal IDs both for the users and the items to make sure that the IDs are consecutive. So to get the items a user interacted with, we have to find out the internal UserID first.

```
>>> internalID=r.getInternalUid("196")
>>> r.getR()[internalID]
set([(521, 1), (377, 4), (365, 3), (438, 5), (86, 4), (649, 4), (0, 3)
    , (522, 3), (423, 3), (389, 5), (751, 3), (656, 4), (947, 4), (432,
    2), (632, 2), (431, 5), (221, 5), (92, 4), (291, 3), (528, 4),
```

---

1 We will use $ to indicate a bash prompt.
2 We will use >>> to indicate the python prompt.

```
    (83, 4), (363, 3), (466, 4), (289, 5), (512, 5), (179, 3), (329, 4)
    , (672, 4), (834, 5), (665, 3), (321, 2), (487, 3), (380, 4),
    (1006, 4), (1045, 3), (491, 3), (302, 4), (550, 5), (10, 2)])
```

`r.getR()` returns a dict with internal UserIDs as keys and sets of (ItemID, NumberOfInteractions) tuples. Please note that the original ID is a string.

To evaluate the algorithms, you have to split the datasets like described at 2.3.1. You do this by calling

```
>>> import util.split
>>> trainingDict, evaluationDict = util.split(r.getR(), 1234567890)
0 Users split.
100 Users split.
200 Users split.
300 Users split.
400 Users split.
500 Users split.
600 Users split.
700 Users split.
800 Users split.
900 Users split.
```

This will split a dict like `r.getR()` returns into a `trainingDict` where one transaction per user is missing and an `evaluationDict` with these missing transactions.

6.3   NON-PERSONALIZED ALGORITHMS

To make our first simple recommendations with the constant recommender, we need to initialize an object of the constant class. As an argument the constructor needs a dict for training:

```
>>> import recommender.nonpersonalized
>>> constant = recommender.nonpersonalized.constant(trainingDict)
>>> constant.getRec(0, 10)
[357, 49, 52, 157, 101, 24, 289, 189, 31, 60]
```

Every recommender has a `getRec` function with this signature. The first argument is the internal UserID, the second is the number of items to be recommended. If n = -1, all items get recommendend. The IDs of the recommended items are internal ones, too. If you want to pass external UserIDs and get external ItemIDs back, you do not have to map them all by yourself. Instead, you can use a helper function called `getExternalRec`.

```
>>> import util.helper
>>> externalConstantgetRec =
    util.helper.getExternalRec(constant.getRec, r)
>>> externalConstantgetRec("196", 10)
['50', '100', '181', '258', '174', '1', '286', '127', '98', '288']
```

You pass the `getRec` function and the reader object and you get a new function which takes and returns only the original IDs.

Now, we want to find out how well this algorithm is performing. Except AUC, the functions for computing the metrics work by letting the recommender recommend for every user as much items as specified in the third argument. With these recommendations and the hidden items, the methods calculate the performance, according to the respective metric. The other arguments are a dict with the hidden items and the getRec function of the respective recommender algorithm.
First we will compute the hitrate:

```
>>> import util.test
>>> util.test.hitrate(evaluationDict, constant.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 4.0
200 users tested
Hits so far: 14.0
300 users tested
Hits so far: 26.0
400 users tested
Hits so far: 31.0
500 users tested
Hits so far: 40.0
600 users tested
Hits so far: 49.0
700 users tested
Hits so far: 52.0
800 users tested
Hits so far: 61.0
900 users tested
Hits so far: 67.0
Number of hits: 69.0
Number of possible hits: 943.0
Hitrate: 0.07317073170731707
```

After some progress output it prints the number of hits and the number of possible hits which is in our case also the number of users. The hitrate is then the number of hits over the number of users. Next, we will calculate the precision:

```
>>> util.test.precision(evaluationDict, constant.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 4.0
200 users tested
Hits so far: 14.0
300 users tested
Hits so far: 26.0
400 users tested
Hits so far: 31.0
500 users tested
Hits so far: 40.0
```

```
600 users tested
Hits so far: 49.0
700 users tested
Hits so far: 52.0
800 users tested
Hits so far: 61.0
900 users tested
Hits so far: 67.0
Number of hits: 69.0
Number of possible hits: 943.0
Precision: 6.9
```

The output looks similar to the one of hitrate and the meaning is analogous. Very similarly we can calculate the F1 metric:

```
>>> util.test.f1(evaluationDict, constant.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 4.0
200 users tested
Hits so far: 14.0
300 users tested
Hits so far: 26.0
400 users tested
Hits so far: 31.0
500 users tested
Hits so far: 40.0
600 users tested
Hits so far: 49.0
700 users tested
Hits so far: 52.0
800 users tested
Hits so far: 61.0
900 users tested
Hits so far: 67.0
Number of hits: 69.0
Number of possible hits: 943.0
F1: 0.14480587618048268
```

Also the mean reciprocal hitrate:

```
>>> util.test.mrhr(evaluationDict, constant.getRec, 10)
0 users tested
Score so far: 0.0
100 users tested
Score so far: 0.675
200 users tested
Score so far: 2.9972222222222227
300 users tested
Score so far: 7.290079365079365
400 users tested
Score so far: 8.34404761904762
500 users tested
Score so far: 10.305158730158729
```

```
600 users tested
Score so far: 12.464682539682537
700 users tested
Score so far: 12.982539682539679
800 users tested
Score so far: 18.332539682539675
900 users tested
Score so far: 19.56825396825396
Score: 19.79325396825396
Number of possible hits: 943.0
Mean Reciprocal Hitrate: 0.02098966486559275
```

To compute the AUC, the third argument has to be our `reader` object because the function needs additional information to compute its metric. We do not need to specify the number of items to recommend, because the AUC works by comparing all items. Thus, it needs the recommender to recommend all available items. The first two arguments do not change.

```
>>> util.test.auc(evaluationDict, constant.getRec, r)
0 users tested
Score so far: 0
100 users tested
Score so far: 77.97155778327935
200 users tested
Score so far: 163.24761858534592
300 users tested
Score so far: 247.0700325949549
400 users tested
Score so far: 329.3602045603418
500 users tested
Score so far: 413.3631291487538
600 users tested
Score so far: 498.4268891693222
700 users tested
Score so far: 581.1334125793506
800 users tested
Score so far: 663.7179944721634
900 users tested
Score so far: 746.0474962705288
AUC: 0.828894041097185
```

These are all evaluation metrics. They work the same way for all recommendation algorithms. You just need to provide the `getRec` function.

The second non-personalized algorithm is the random recommender. You can use it like this:

```
>>> randomRec = recommender.nonpersonalized.randomRec(trainingDict,
    1234567890)
>>> randomRec.getRec(0, 10)
[1548, 1068, 746, 1357, 1501, 1359, 428, 141, 233, 726]
>>> util.test.hitrate(evaluationDict, randomRec.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
```

```
Hits so far: 2.0
200 users tested
Hits so far: 3.0
300 users tested
Hits so far: 3.0
400 users tested
Hits so far: 4.0
500 users tested
Hits so far: 4.0
600 users tested
Hits so far: 4.0
700 users tested
Hits so far: 4.0
800 users tested
Hits so far: 5.0
900 users tested
Hits so far: 5.0
Number of hits: 5.0
Number of possible hits: 943.0
Hitrate: 0.005302226935312832
```

The second argument of the constructor is the seed for the random function.

From now on, we will only show the hitrate, because the other metrics are computed similarly. But feel free to also compute the other metrics as well.

6.4 k-nearest neighbor

Now we want to try the first more sophisticated algorithm, namely the item based k-Nearest Neighbor. The k-Nearest Neighbor algorithms need the database as a matrix so we have to first split the matrix, provided by the reader object.

```
>>> trainingMatrix, matrixEvaluationDict = util.split.splitMatrix(r.
    getMatrix, 123456789)
0 Users split.
100 Users split.
200 Users split.
300 Users split.
400 Users split.
500 Users split.
600 Users split.
700 Users split.
800 Users split.
900 Users split.
```

Depending on the recommendation algorithm, we need a matrix or a dict. So here we pass a matrix like r.getMatrix() returns and get a trainingMatrix, where one entry per user is set to zero and a dict with these missing entries. To understand the matrix represention of the dataset refer to 4.2.1.

Thus we initialize an itemKnn object. The initialization builds the model yet, so it will need some time until the constructor has finished.

```
>>> import recommender.knn
>>> itemKnn = recommender.knn.itemKnn(trainingMatrix, 10)
0 Similarities calculated
10000 Similarities calculated
20000 Similarities calculated
30000 Similarities calculated
...
1400000 Similarities calculated
1410000 Similarities calculated
```

Please note that you pass the the right dict for evaluation as the first argument. In this case it is `matrixEvaluationDict`. The second argument is the number of neighbors to include into the neighborhood. During the model building, phase the algorithm will compute the similarity of each item to each other item. But because this similarity is symmetric, it is only necessary to compute the half. The 100k movieLens dataset we use for these examples has 1682 items, so we have to compute $\frac{N(N-1)}{2} = \frac{1682*1681}{2} = 1413721$ similarities.

Now you can test this algorithm with the test metrics for example the hitrate:

```
>>> util.test.hitrate(matrixEvaluationDict, itemKnn.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 21.0
200 users tested
Hits so far: 54.0
300 users tested
Hits so far: 78.0
400 users tested
Hits so far: 97.0
500 users tested
Hits so far: 126.0
600 users tested
Hits so far: 152.0
700 users tested
Hits so far: 181.0
800 users tested
Hits so far: 211.0
900 users tested
Hits so far: 239.0
Number of hits: 248.0
Number of possible hits: 943.0
Hitrate: 0.26299045599151644
```

As you can see, this algorithm achieves a much better performance than the non-personalized algorithms.

For user based k-Nearest Neighbors, the call to build the model looks like this:

```
>>> userKnn = recommender.knn.userKnn(trainingMatrix, 50)
```

```
0 Similarities calculated
10000 Similarities calculated
20000 Similarities calculated
30000 Similarities calculated
...
430000 Similarities calculated
440000 Similarities calculated
```

The dataset has 943 users, so this time we have to compute $\frac{N(N-1)}{2} = \frac{943*942}{2} = 444153$ similarities. To get the hitrate for this algorithm, we call the hitrate function with the getRec function just like before.

```
>>> util.test.hitrate(matrixEvaluationDict, userKnn.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 21.0
200 users tested
Hits so far: 52.0
300 users tested
Hits so far: 82.0
400 users tested
Hits so far: 102.0
500 users tested
Hits so far: 130.0
600 users tested
Hits so far: 153.0
700 users tested
Hits so far: 173.0
800 users tested
Hits so far: 206.0
900 users tested
Hits so far: 237.0
Number of hits: 247.0
Number of possible hits: 943.0
Hitrate: 0.26193001060445387
0.26193001060445387
```

6.5 BPRMF

For matrix factorization algorithms like BPRMF, RankMFX and Ranking SVD, you need to specify some meta parameters.

```
>>> import recommender.BPRMF
>>> W, H = recommender.BPRMF.learnModel(r.getMaxUid(), r.getMaxIid(),
                        0.01, 0.01, 0.01,   # regularization parameter
                        0.1,                # learning rate
                        trainingDict,       # training dict
                        150,                # number of features
                        3,                  # number of epochs
                        r.numberOfTransactions)
Epoch: 1/3 | iteration 1000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.697530
```

```
Epoch: 1/3 | iteration 2000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.694108
Epoch: 1/3 | iteration 3000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.695815
...
Epoch: 3/3 | iteration 99000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.124050
Epoch: 3/3 | iteration 100000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.125070
```

This will output quite much, so you can observe how the loss is behaving. With a short search, we found these values for the regularization parameters and the learning rate to be quite good, but they are probably not the best. But above all, the results will get better with more epochs, but then it will, of course, take longer to compute. Thus, three is alright for a short experiment. The `learnModel` returns a model in the form of two matrices. Two get a getRec method, you have to instantiate a `Mfrec` class from the `mf` module in `recommender.mf`, which offers such a method.

```
>>> BPRMF = recommender.mf.MFrec(W, H, trainingDict)
>>> util.test.hitrate(evaluationDict, BPRMF.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 26.0
200 users tested
Hits so far: 51.0
300 users tested
Hits so far: 78.0
400 users tested
Hits so far: 95.0
500 users tested
Hits so far: 122.0
600 users tested
Hits so far: 148.0
700 users tested
Hits so far: 171.0
800 users tested
Hits so far: 203.0
900 users tested
Hits so far: 224.0
Number of hits: 235.0
Number of possible hits: 943.0
Hitrate: 0.2492046659597031
```

As noted above, the results get better with more epochs.

6.6 RANKMFX

```
>>> import recommender.RankMFX
>>> W, H = recommender.RankMFX.learnModel(r.getMaxUid(),r.getMaxIid(),
                    0.01, 0.01, 0.01, # regularization parameter
```

```
                   0.1,              # learning rate
                   trainingDict,     # training dict
                   250,              # number of features
                   5,                # number of epochs
                   r.numberOfTransactions)
Epoch: 1/5 | iteration 1000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.984394
Epoch: 1/5 | iteration 2000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.978983
Epoch: 1/5 | iteration 3000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.966176
...
Epoch: 5/5 | iteration 99000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.141662
Epoch: 5/5 | iteration 100000/100000 | learning rate=0.100000 |
    average_loss for the last 1000 iterations = 0.130552
```

Again, there is much output to monitor the algorithm. As above, we have to generate a `getRec` method before we can evaluate the algorithm.

```
>>> RankMFX = recommender.mf.MFrec(W, H, trainingDict)
>>> util.test.hitrate(evaluationDict, RankMFX.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 20.0
200 users tested
Hits so far: 37.0
300 users tested
Hits so far: 54.0
400 users tested
Hits so far: 73.0
500 users tested
Hits so far: 103.0
600 users tested
Hits so far: 123.0
700 users tested
Hits so far: 136.0
800 users tested
Hits so far: 154.0
900 users tested
Hits so far: 172.0
Number of hits: 177.0
Number of possible hits: 943.0
Hitrate: 0.18769883351007424
```

What applied to BPRMF in terms of performance also applies to RankMFX: With better tuned meta arguments and more time, the results will likely be better.

6.7   ranking svd (sparse svd)

Next, the Ranking SVD or Sparse SVD algorithm: To build the model, we have a method similar to RankMFX and BPRMF, just without regularization parameters.

```
>>> import recommender.svd
>>> W, H = recommender.svd.learnModel(r.getMaxUid(), r.getMaxIid(),
                    0.0002,         # learning rate
                    trainingDict,   # training dict
                    770,            # number of features
                    40,             # number of epochs
                    1000)           # number of iterations
Epoch: 1/40 | iteration 100/1000 | learning rate=0.000200 |
    average_loss for the last 100 iterations = -3.451531 | time needed:
     0.420000
Epoch: 1/40 | iteration 200/1000 | learning rate=0.000200 |
    average_loss for the last 100 iterations = -3.687917 | time needed:
     0.400000
Epoch: 1/40 | iteration 300/1000 | learning rate=0.000200 |
    average_loss for the last 100 iterations = -3.433482 | time needed:
     0.410000
...
Epoch: 40/40 | iteration 900/1000 | learning rate=0.000200 |
    average_loss for the last 100 iterations = -2.934971 | time needed:
     0.410000
Epoch: 40/40 | iteration 1000/1000 | learning rate=0.000200 |
    average_loss for the last 100 iterations = -2.734519 | time needed:
     0.400000
```

The evaluation works the same way, as for the other matrix factorization algorithms:

```
>>> svd = recommender.mf.MFrec(W, H, trainingDict)
>>> util.test.hitrate(evaluationDict, svd.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 7.0
200 users tested
Hits so far: 19.0
300 users tested
Hits so far: 29.0
400 users tested
Hits so far: 37.0
500 users tested
Hits so far: 54.0
600 users tested
Hits so far: 71.0
700 users tested
Hits so far: 77.0
800 users tested
Hits so far: 95.0
900 users tested
```

```
Hits so far: 108.0
Number of hits: 113.0
Number of possible hits: 943.0
Hitrate: 0.11983032873807
```

At last the slope one recommender:

```
>>> import recommender.slopeone
>>> slopeone = recommender.slopeone.slopeone(trainingDict)
100000 differences computed
200000 differences computed
300000 differences computed
...
19800000 differences computed
19900000 differences computed
```

And its evaluation:

```
>>> util.test.hitrate(evaluationDict, slopeone.getRec, 10)
0 users tested
Hits so far: 0.0
100 users tested
Hits so far: 11.0
200 users tested
Hits so far: 21.0
300 users tested
Hits so far: 28.0
400 users tested
Hits so far: 38.0
500 users tested
Hits so far: 49.0
600 users tested
Hits so far: 61.0
700 users tested
Hits so far: 70.0
800 users tested
Hits so far: 76.0
900 users tested
Hits so far: 78.0
Number of hits: 81.0
Number of possible hits: 943.0
Hitrate: 0.08589607635206786
```

CONCLUSIONS

In this thesis, we provided an overview over the research area by explaining the terminology of recommender systems, respectively collaborative filtering and their task. We also introduced the leave-one-out protocol used to evaluate the performance of recommender systems.

This is followed by explanations of the recommender algorithms implemented in recsyslab together with the core part of their implementation. By comparing the pseudocode and the Python code implementing it, we can see, that the implementation is quite straightforward.

After we provided the necessary requirements, we showed how the described algorithms perform with the also described metrics and we showed a high level view on recsyslab to help developers extend it.

All this should provide enough knowledge for the reader to understand roughly how recsyslab works in theory. So we finally got to the extensive user manual, explaining how to use every recommender algorithm and test metric, so the reader would be able to use the library.

The user manual shows that it is really easy to use recsyslab. Together with the comparisons of the Python code with the pseudocode in Chapter 4, we see that the goals of this thesis are achieved.

Recsyslab provides several state of the art recommender algorithms, several widely used test metrics, and a simple infrastructure to use, train and test recommender algorithms. Additionally, recsyslab is easy extendable with new recommender algorithms or test metrics and it produces acceptable test results. After reading this document, every beginner should be able to use and understand recsyslab.

We hope this will help students start working in this field and researchers to compare their new algorithms with already existing algorithms without wasting much time implementing something else then their own algorithms. This was the main motivation to write this thesis.

7.1 OUTLOOK

For future work, we could implement more recommender algorithms, implement more test metrics or visualize test results graphically. It would also be useful to update the model with new interactions, so one does not have to train the model from ground up again, or to have the possibility to continue training models after they got evaluated.

We plan to continue testing our algorithms to optimize its performance and to detect any possible bugs. Furthermore, we plan to keep the project active, and to engage more participants, so as to conduct a more comprehensive and thorough testing of the library.

We also plan to integrate functionality to support grid-search for hyperparameter tuning.

Apart from that, we hope to get feedback from users to improve recsyslab. Thus, it would be interesting to see, which parts of the implementation students might have problems with, when recsyslab is used to support a recommender system lecture. Then we could refine these parts. Furthermore, we expect to receive direct feedback from practitioners regarding which new features, test metrics and additional recommender algorithms should be implemented to keep recsyslab up to date.

APPENDIX - DOCSTRINGS

In this chapter, we provide the inline documentation which is also included in recsyslab.

```
Help on package bin:

NAME
    bin - Package with the Python scripts of the recsyslab toolbox.

FILE
    /home/jk/recsyslab/bin/__init__.py

DESCRIPTION
    It contains the Packages:
        recommender -- Package for the recommender algorithms
        util        -- Package with all the other stuff needed to try out the
                        recommender algorithms

PACKAGE CONTENTS
    experiments
    main
    manual
    recommender (package)
    testing
    usecases
    util (package)
```

```
Help on package bin.recommender in bin:

NAME
    bin.recommender - Package containing the recommending algorithms.

FILE
    /home/jk/recsyslab/bin/recommender/__init__.py

DESCRIPTION
    These are:
        nonpersonalized --  constant and random
        knn             --  k-Nearest-Neigbor
        mf              --  Matrix Factorization with an arbitrary loss
        BPRMF           --  Bayesian Personalized Ranking with MF
                            (just calls mf with logLoss)
        RankMFX         --  MF with hingeLoss
        slopeone        --  slopeone recommender
        svd             --  Ranking SVD (Sparse SVD)

PACKAGE CONTENTS
    BPRMF
    RankMFX
    knn
    mf
    nonpersonalized
    slopeone
    svd
```

```
Help on module bin.recommender.knn in bin.recommender:

NAME
    bin.recommender.knn - Two classes for k-Nearest-Neighbor(knn) recommendation.

FILE
    /home/jk/recsyslab/bin/recommender/knn.py

DESCRIPTION
    itemKnn --  Item based knn
    userKnn --  User based knn

    Based on:
    George Karypis. 2001. Evaluation of Item-Based Top-N Recommendation
    Algorithms. In Proceedings of the tenth international conference on
    Information and knowledge management (CIKM 01), Henrique Paques,
    Ling Liu, and David Grossman (Eds.). ACM, New York, NY, USA, 247-254.
    DOI=10.1145/502585.502627 http://doi.acm.org/10.1145/502585.502627

CLASSES
    __builtin__.object
        itemKnn
```

```
        userKnn

    class itemKnn(__builtin__.object)
     |  Class for item based knn.
     |
     |  Methods defined here:
     |
     |  __init__(self, userItemMatrix, n)
     |      Builds a model for Item based knn.
     |
     |          userItemMatrix  --  A matrix where an entry at i, j is the rating
     |                              the ith user gave the jth item.
     |          n               --  number of neighbors which are getting computed.
     |
     |      Uses the cosine for similarity.
     |
     |  getRec(self, u, n)
     |      Returns the n best recommendations for user u.
     |
     |      Set n = -1 to get all items recommended
     |
     |  ----------------------------------------------------------
     |  Data descriptors defined here:
     |
     |  __dict__
     |      dictionary for instance variables (if defined)
     |
     |  __weakref__
     |      list of weak references to the object (if defined)

    class userKnn(__builtin__.object)
     |  Methods defined here:
     |
     |  __init__(self, userItemMatrix, n)
     |      Builds a model for user based knn.
     |
     |          userItemMatrix  --  A matrix where an entry at i, j is the rating
     |                              the ith user gave the jth item.
     |          n               --  number of neighbors which get computed.
     |
     |      Uses the cosine for similarity.
     |
     |  getRec(self, u, n)
     |      Returns the n best recommendations for user u.
     |
     |      Set n = -1 to get all items recommended
     |
     |  ----------------------------------------------------------
     |  Data descriptors defined here:
     |
     |  __dict__
     |      dictionary for instance variables (if defined)
     |
     |  __weakref__
     |      list of weak references to the object (if defined)

FUNCTIONS
    computeCosSim(sim, matrix)
        Computes the cos of every two lines in the matrix and writes them into sim.

    cos(a, b)
        Gets two vectors(onedimensional np.matrix) and computes their cosine.
```

```
Help on module bin.recommender.mf in bin.recommender:

NAME
    bin.recommender.mf - Module for Matrix Factorization (mf) techniques.

FILE
    /home/jk/recsyslab/bin/recommender/mf.py

DESCRIPTION
    learnModel  -- learn a mf model
    MFrec       -- compute recommendations based on a mf model

CLASSES
    __builtin__.object
        MFrec

    class MFrec(__builtin__.object)
     |  Class to compute recommendations with a mf model.
     |
     |  Methods defined here:
     |
     |  __init__(self, W, H, trainingR)
     |      Initialize the class.
     |
     |      W           -- Matrix of the User features
     |      H           -- Matrix of the Item features
     |      trainingR   -- The dict the model was trained with
     |
     |  getRec(self, u, n)
     |      Returns the n best recommendations for user u based on the mf model.
     |
     |      Set n = -1 to get all items recommended
     |
     |  ----------------------------------------------------------
     |  Data descriptors defined here:
     |
     |  __dict__
     |      dictionary for instance variables (if defined)
     |
     |  __weakref__
     |      list of weak references to the object (if defined)

FUNCTIONS
    learnModel(n_users, m_items, regU, regI, regJ, learningRate, R, features, epochs, numberOfIterations, lossF,
            dlossF)
        Learns a mf model with a passed loss function.

            n_users           --  The highest internal assigned User ID
            m_items           --  The highest internal assigned Item ID
            regU              --  Regularization for the user vector
            regI              --  Regularization for the positive item
            regJ              --  Regularization for the negative item
            learningRate      --  The learning rate
            R                 --  A dict of the form UserID -> (ItemId, Rating)
            features          --  Number of features of the items and users
            epochs            --  Number of epochs the model should be learned
            numberOfIterations --  Number of iterations in each epoch
            lossF             --  Loss function
            dlossF            --  Derivation of lossF

        Returns:
            W  --  User Features
            H  --  Item Features

        Based on:
        Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and
        Lars Schmidt-Thieme. 2009. BPR: Bayesian personalized ranking from
        implicit feedback. In Proceedings of the Twenty-Fifth Conference on
        Uncertainty in Artificial Intelligence (UAI 09).
        AUAI Press, Arlington, Virginia, United States, 452-461.
```

```
Help on module bin.recommender.BPRMF in bin.recommender:

NAME
    bin.recommender.BPRMF

FILE
    /home/jk/recsyslab/bin/recommender/BPRMF.py

FUNCTIONS
    dLogLoss(a, y)
        Derivative of the logLoss.

    learnModel(n_users, m_items, regU, regI, regJ, learningRate, R, features, epochs, numberOfIterations)
        Learns a model with the BPRMF algorithm, actually just calls mf.

            n_users           --  The highest internal assigned User ID
            m_items           --  The highest internal assigned Item ID
            regU              --  Regularization for the user vector
```

```
                  regI              -- Regularization for the positive item
                  regJ              -- Regularization for the negative item
                  learningRate      -- The learning rate
                  R                 -- A dict of the form UserID -> (ItemId, Rating)
                  features          -- Number of features of the items and users
                  epochs            -- Number of epochs the model should be learned
                  numberOfIterations -- Number of iterations in each epoch

          Returns:
              W   --  User Features
              H   --  Item Features

          Based on:
          Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and
          Lars Schmidt-Thieme. 2009. BPR: Bayesian personalized ranking from
          implicit feedback. In Proceedings of the Twenty-Fifth Conference on
          Uncertainty in Artificial Intelligence (UAI 09).
          AUAI Press, Arlington, Virginia, United States, 452-461.

      logLoss(a, y)
          logLoss(a, y) = log(1 + exp(-a*y))
```

```
Help on module bin.recommender.RankMFX in bin.recommender:

NAME
    bin.recommender.RankMFX

FILE
    /home/jk/recsyslab/bin/recommender/RankMFX.py

FUNCTIONS
    dHingeLoss(a, y)
        Derivative of the hingeLoss.

    hingeLoss(a, y)
        hingeLoss(a, y) = max(0, 1 - a*y)

    learnModel(n_users, m_items, regU, regI, regJ, learningRate, R, features, epochs, numberOfIterations)
        Learns a model with the RankMFX algorithm, actually just calls mf.

              n_users           -- The highest internal assigned User ID
              m_items           -- The highest internal assigned Item ID
              regU              -- Regularization for the user vector
              regI              -- Regularization for the positive item
              regJ              -- Regularization for the negative item
              learningRate      -- The learning rate
              R                 -- A dict of the form UserID -> (ItemId, Rating)
              features          -- Number of features of the items and users
              epochs            -- Number of epochs the model should be learned
              numberOfIterations -- Number of iterations in each epoch

          Returns:
              W   --  User Features
              H   --  Item Features
```

```
Help on module bin.recommender.slopeone in bin.recommender:

NAME
    bin.recommender.slopeone

FILE
    /home/jk/recsyslab/bin/recommender/slopeone.py

CLASSES
    __builtin__.object
        slopeone

    class slopeone(__builtin__.object)
     |  A class to compute recommendations with a Slope One Predictor.
     |
     |  Based on:
     |      "Slope One Predictors for Online Rating-Based Collaborative Filtering"
     |      by Daniel Lemire and Anna Maclachlan.
     |
     |  Methods defined here:
     |
     |  __init__(self, R)
     |      Generate the model.
     |
     |      R   --  A dict of the form UserID -> (ItemId, Rating)
     |
     |  getRec(self, u, n)
     |      Returns the n best recommendations for user u.
     |
     |      Set n = -1 to get all items recommended
     |
     |  ----------------------------------------------------------
     |  Data descriptors defined here:
```

```
    |
    | __dict__
    |     dictionary for instance variables (if defined)
    |
    | __weakref__
    |     list of weak references to the object (if defined)
```

```
Help on module bin.recommender.svd in bin.recommender:

NAME
    bin.recommender.svd - Module to compute the model of Ranking SVD (Sparse SVD)

FILE
    /home/jk/recsyslab/bin/recommender/svd.py

DESCRIPTION
    Based on:
    Jahrer, Michael, and Andreas Toescher.
    "Collaborative filtering ensemble for ranking."
    Proc. of KDD Cup Workshop at 17th ACM SIGKDD Int. Conf.
    on Knowledge Discovery and Data Mining, KDD. Vol. 11. 2011.

FUNCTIONS
    learnModel(n_users, m_items, learningRate, R, features, epochs, numberOfIterations)
        Returns the model of a learned svd as two matrices.

        n_users            --  The highest internal assigned User ID
        m_items            --  The highest internal assigned Item ID
        learningRate       --  The learning rate
        R                  --  A dict of the form UserID -> (ItemId, Rating)
        features           --  Number of features of the items and users
        epochs             --  Number of epochs the model should be learned
        numberOfIterations --  Number of iterations in each epoch
```

```
Help on package bin.util in bin:

NAME
    bin.util - Package with utility code to run the recommender algorithms.

FILE
    /home/jk/recsyslab/bin/util/__init__.py

DESCRIPTION
    It contains the following modules:
        reader  --  A reader to read in databases
        split   --  Split databases
        test    --  Test metrics
        helper  --  Several helping methods and a two way dict

PACKAGE CONTENTS
    helper
    reader
    split
    test
```

```
Help on module bin.util.reader in bin.util:

NAME
    bin.util.reader - Contains a class to manage a database.

FILE
    /home/jk/recsyslab/bin/util/reader.py

CLASSES
    __builtin__.object
        stringSepReader

    class stringSepReader(__builtin__.object)
     |  A class to manage a database.
     |
     |  Methods defined here:
     |
     |  __init__(self, filename, separator)
     |      Reads in a database file.
     |
     |      The lines of the database file have to look like this:
     |          UserID<separator>ItemID<separator>Rating
     |      If there is just an UserID and an ItemID the rating is set to 1.
     |      Everything coming after the rating will be ignored,
     |      but when omit the rating and still have something following the ItemID
     |      this will be understood as the rating.
     |
     |  getInternalIid(self, originalIid)
     |      Maps the given original ItemID to the corresponding internal ItemID.
     |
     |      The passed original UserID has to be a string.
     |
```

```
      |  getInternalUid(self, originalUid)
      |      Maps the given original UserID to the corresponding internal UserID.
      |
      |      The passed original UserID has to be a string.
      |
      |  getMatrix(self)
      |      Get the database as a matrix.
      |
      |      The lines of the matrix are corresponding to the users
      |      and the columns to the items.
      |      So the n,m entry is the rating the nth user gave the mth item.
      |
      |  getMaxIid(self)
      |      Returns the highest internal assigned ItemID.
      |
      |  getMaxUid(self)
      |      Returns the highest internal assigned UserID.
      |
      |  getOriginalIid(self, internalIid)
      |      Maps the given internal ItemID to the corresponding original ItemID.
      |
      |  getOriginalUid(self, internalUid)
      |      Maps the given internal UserID to the corresponding original UserID.
      |
      |  getR(self)
      |      Return the database as a dict.
      |
      |      The dict has internal UserIDs as keys and
      |      (ItemID, Rating) Tuples as values
      |
      |  ----------------------------------------------------------
      |  Data descriptors defined here:
      |
      |  __dict__
      |      dictionary for instance variables (if defined)
      |
      |  __weakref__
      |      list of weak references to the object (if defined)
```

```
Help on module bin.util.split in bin.util:

NAME
    bin.util.split - Methods to split a database with the leave-one-out method.

FILE
    /home/jk/recsyslab/bin/util/split.py

DESCRIPTION
    split        --  splits a reader into two new dicts
    splitMatrix  --  splits a matrix into a matrix and a dict

FUNCTIONS
    split(R, seed)
        Splits a database into two dicts.

        Splits after the leave-one-out method which means for every user in the
        database it takes one item out of the database and into a new one.
        The first returned dict is the database with one item missing for each
        user. The second returned dict has the missing items.

        R is a dict of the database.
        seed is a seed for the randomness.

    splitMatrix(M, seed)
        Splits a matrix into two new dicts.

        Returns an User x Item Matrix with one entry of each user set to 0
        and a User -> Item Dict with the missing entrys.

        M is an User x Item Matrix.
        seed is a seed for the randomness.
```

```
Help on module bin.util.test in bin.util:

NAME
    bin.util.test - Module with several metrics to test the recommender algorithms.

FILE
    /home/jk/recsyslab/bin/util/test.py

DESCRIPTION
    hitrate     --  Returns #hits, #items
    precision   --  Returns #hits / n
    f1          --  Returns the result of a F1 test
    mrhr        --  Returns the Mean Reciprocal Hitrate
    auc         --  Returns the Area under the curve
    countHits   --  Returns hits, items in a tuple
```

```
FUNCTIONS
    auc(testR, recommender, r)
        Returns and prints the AUC of the recommender function.

        testR is a dict with an internal UserID as a dict and a list of items as
        values. Normally testR is the second dict split.split returns.
        The list can have a lenght greater than 1.

        recommender is a function which takes an internal UserID and n and returns
        n items recommender for the UserID.

        r is a reader object with the database read in.

        Based on:
        Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and
        Lars Schmidt-Thieme. 2009. BPR: Bayesian personalized ranking from
        implicit feedback. In Proceedings of the Twenty-Fifth Conference on
        Uncertainty in Artificial Intelligence (UAI 09).
        AUAI Press, Arlington, Virginia, United States, 452-461.

    countHits(testR, recommender, n)
        Returns the number of hits and the number of items it was tested with.

        Is a helper function for the other metrics
        testR is a dict with an internal UserID as a dict and a list of items as
        values. Normally testR is the second dict split.split returns.
        The list can have a lenght greater than 1.

        recommender is a function which takes an internal UserID and n and returns
        n items recommender for the UserID.

        n is the number of items the recommender can recommend.

        hits = number of items from testR the recommender guessed right.
        items = the number of items in testR.

        hits, items will be printed.
        hits, items will get returned.

    f1(testR, recommender, n)
        Prints and returns F1 of the recommender.

        F1 = (2 * hitrate * precision) / (hitrate + precision)
        testR is a dict with an internal UserID as a dict and a list of items as
        values. Normally testR is the second dict split.split returns.
        The list can have a lenght greater than 1.

        recommender is a function which takes an internal UserID and n and returns
        n items recommender for the UserID.

        n is the number of items the recommender can recommend.

        hits = number of items from testR the recommender guessed right.

        See also: "Application of Dimensionality Reduction in Recommender System
        -- A Case Study" by Badrul M. sarcar et al.

    hitrate(testR, recommender, n)
        Returns the number of hits and the number of items it was tested with.

        testR is a dict with an internal UserID as a dict and a list of items as
        values. Normally testR is the second dict split.split returns.
        The list can have a lenght greater than 1.

        recommender is a function which takes an internal UserID and n and returns
        n items recommender for the UserID.

        n is the number of items the recommender can recommend.

        hits = number of items from testR the recommender guessed right.
        items = the number of items in testR.

        hits, items and hits / items will be printed.
        hits, items will get returned.

        hits / items gives the hitrate or recall.

        See also:
        Sarwar, Badrul, et al. Application of dimensionality reduction in
        recommender system-a case study. No. TR-00-043.
        MINNESOTA UNIV MINNEAPOLIS DEPT OF COMPUTER SCIENCE, 2000.

    mrhr(testR, recommender, n)
        Returns the Mean Reciprocal Hitrate of the recommender.

        testR is a dict with an internal UserID as a dict and a list of items as
        values. Normally testR is the second dict split.split returns.
        The list can have a length greater than 1.

        recommender is a function which takes an internal UserID and n and returns
        n items recommender for the UserID.
```

        n is the number of items the recommender can recommend.

        hits = number of items from testR the recommender guessed right.

    precision(testR, recommender, n)
        Returns the number of hits / n .

        testR is a dict with an internal UserID as a dict and a list of items as
        values. Normally testR is the second dict split.split returns.
        The list can have a lenght greater than 1.

        recommender is a function which takes an internal UserID and n and returns
        n items recommender for the UserID.

        n is the number of items the recommender can recommend.

        hits = number of items from testR the recommender guessed right.

        See also:
        Sarwar, Badrul, et al. Application of dimensionality reduction in
        recommender system-a case study. No. TR-00-043.
        MINNESOTA UNIV MINNEAPOLIS DEPT OF COMPUTER SCIENCE, 2000.

Help on module bin.util.helper in bin.util:

NAME
    bin.util.helper - Several helper functions and a helper class.

FILE
    /home/jk/recsyslab/bin/util/helper.py

DESCRIPTION
    idOrigDict   -- two way dict
    normRowNorm   -- normalize a matrix by the norm
    normRowSum   -- normalize a matrix by the sum
    cache   -- cache results of a function
    writeOrigToFile   -- save a DB with original IDs
    writeInternalToFile -- save a DB with internal IDs
    sortList   -- sort a list with scores and IDs
    listToSet   -- convert a list into a set
    sortResults   -- sort test results
    printMatrix   -- print a matrix
    randDataset   -- generate a random DB
    dictToMatrix   -- convert a dict to a matrix
    getScoreMF   -- compute the score of an item
    getExternalRec   -- convert getRec for original IDs

CLASSES
    __builtin__.object
        idOrigDict

    class idOrigDict(__builtin__.object)
     | Class two map external/original IDs to internal IDs and vice versa.
     |
     | Methods defined here:
     |
     | __init__(self)
     | Initialization.
     |
     | add(self, orig)
     | Add a new original ID.
     |
     | Returns the internal ID the passed original ID got mapped to.
     | If the passed ID is already mapped nothing happens except the
     | mapped internal ID is returned.
     |
     | getId(self, orig)
     | Returns the internal ID which is mapped to the passed orinigal ID.
     |
     | getOrig(self, id)
     | Returns the original ID which is mapped to the passed internal ID.
     |
     | --
     | Data descriptors defined here:
     |
     | __dict__
     | dictionary for instance variables (if defined)
     |
     | __weakref__
     | list of weak references to the object (if defined)

FUNCTIONS
    cache(fn)
        Decorator to cache the result of the function fn.

    dictToMatrix(d)
        Converts a dict like reader.getR() returns to a matrix.

    getExternalRec(getRec, r)

```
        Converts getRec so it takes and returns only the original IDs.

        r   --  reader object

getScoreMF(origUid, origIid, W, H, r)
        Returns the score of one item for one user with an MF model.
        origUid --  original User ID
        origIid --  original Item ID
        W, H    --  The MF model like returned by recommender.mf for example
        r       --  a reader object with the database

listToSet(sortedscorelist, n)
        Converts a list into a set with size n.

        Gets a list like sortList returns and returns a set with the first n
        items, or less, when there are not enough items

normRowNorm(m)
        Normalize the matrix in place so each row of the matrix has norm 1.

normRowSum(m)
        Normalize the matrix in place so each row of the matrix has sum 1.

printMatrix(m)
        Prints a whole matrix.

        Prints the whole matrix m without replacing stuff with dots
        like it normally does when the matrix is to large.

randDataset(user, items, p, seed, filename)
        Generates a random dataset and writes it into a file.

        user    --  number of users
        items   --  number of items
        p       --  percentage of ones in the dataset
        seed    --  seed for the randomness

sortList(scorelist)
        Gets a list of tuples (itemid, score), sorts by score decreasing.

sortResults(name)
        Sort test results written to a file.

writeInternalToFile(reader, toSave, filename)
        Writes a database into a file with the internal IDs.

            reader  --  reader object
            toSave  --  Part of the database to write
            filename--  Name of the file

        Writes toSave into a file with the name filename.

writeOrigToFile(reader, toSave, filename)
        Writes a database into a file with the original IDs.

            reader  --  reader object
            toSave  --  Part of the database to write
            filename--  Name of the file

        Writes toSave into a file with the name filename but first the
        internal IDs in toSave get mapped to the original IDs.
```

BIBLIOGRAPHY

[1] Gnu general public license, version 3. http://www.gnu.org/licenses/gpl.html, June 2007. Last retrieved 2013-06-10.

[2] Cofirank, June 2013. URL https://sites.google.com/a/cofirank.org/index/Home.

[3] Collaborative filtering, June 2013. URL http://en.wikipedia.org/wiki/Collaborative_filtering.

[4] Content-based filtering, June 2013. URL http://en.wikipedia.org/wiki/Recommender_system#Content-based_filtering.

[5] Duine framework, June 2013. URL http://www.duineframework.org/.

[6] The leave-one-out protocol, June 2013. URL http://en.wikipedia.org/wiki/Cross-validation_%28statistics%29#Leave-one-out_cross-validation.

[7] Apache mahout, June 2013. URL http://mahout.apache.org/.

[8] Matrix factorization, June 2013. URL http://en.wikipedia.org/wiki/Matrix_factorization.

[9] Movielens data sets, June 2013. URL http://grouplens.org/node/73.

[10] Netflix prize, June 2013. URL http://www.netflixprize.com/.

[11] Numpy, June 2013. URL http://www.numpy.org/.

[12] Python programming language, June 2013. URL http://www.python.org/.

[13] recsyslab, June 2013. URL https://github.com/Foolius/recsyslab.

[14] User based knn, June 2013. URL https://github.com/zenogantner/MyMediaLite/blob/master/src/MyMediaLite/RatingPrediction/UserKNN.cs.

[15] Thierry Bertin-Mahieux, Daniel P.W. Ellis, Brian Whitman, and Paul Lamere. The million song dataset. In *Proceedings of the 12th International Conference on Music Information Retrieval (ISMIR 2011)*, 2011.

[16] John S Breese, David Heckerman, and Carl Kadie. Empirical analysis of predictive algorithms for collaborative filtering. In *Proceedings of the Fourteenth conference on Uncertainty in artificial intelligence*, pages 43–52. Morgan Kaufmann Publishers Inc., 1998.

[17] Ernesto Diaz-Aviles, Lucas Drumond, Zeno Gantner, Lars Schmidt-Thieme, and Wolfgang Nejdl. What is happening right now... that interests me?: online topic discovery and recommendation in twitter. In *Proceedings of the 21st ACM international conference on Information and knowledge management*, pages 1592–1596. ACM, 2012.

[18] Michael D. Ekstrand, Michael Ludwig, Joseph A. Konstan, and John T. Riedl. Rethinking the recommender research ecosystem: reproducibility, openness, and lenskit. In *Proceedings of the fifth ACM conference on Recommender systems*, RecSys '11, pages 133–140, New York, NY, USA, 2011. ACM. ISBN 978-1-4503-0683-6. doi: 10.1145/2043932.2043958. URL http://doi.acm.org/10.1145/2043932.2043958.

[19] Zeno Gantner, Steffen Rendle, Christoph Freudenthaler, and Lars Schmidt-Thieme. MyMediaLite: A free recommender system library. In *Proceedings of the 5th ACM Conference on Recommender Systems (RecSys 2011)*, 2011.

[20] Michael Jahrer and Andreas Töscher. Collaborative filtering ensemble for ranking. In *Proc. of KDD Cup Workshop at 17th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining, KDD*, volume 11, 2011.

[21] George Karypis. Evaluation of item-based top-n recommendation algorithms. In *Proceedings of the tenth international conference on Information and knowledge management*, CIKM '01, pages 247–254, New York, NY, USA, 2001. ACM. ISBN 1-58113-436-3. doi: 10.1145/502585.502627. URL http://doi.acm.org/10.1145/502585.502627.

[22] Neil D Lawrence and Raquel Urtasun. Non-linear matrix factorization with gaussian processes. In *Proceedings of the 26th Annual International Conference on Machine Learning*, pages 601–608. ACM, 2009.

[23] J. Lee, M. Sun, and G. Lebanon. A Comparative Study of Collaborative Filtering Algorithms. *ArXiv e-prints*, May 2012.

[24] Daniel Lemire and Anna Maclachlan. Slope one predictors for online rating-based collaborative filtering. *Society for Industrial Mathematics*, 5: 471–480, 2005.

[25] Daniel Lemire and Anna Maclachlan. Slope one predictors for online rating-based collaborative filtering. *CoRR*, abs/cs/0702144, 2007.

[26] Greg Linden, Brent Smith, and Jeremy York. Amazon. com recommendations: Item-to-item collaborative filtering. *Internet Computing, IEEE*, 7(1):76–80, 2003.

[27] Xia Ning and George Karypis. Slim: Sparse linear methods for top-n recommender systems. In Diane J. Cook, Jian Pei, Wei Wang, Osmar R. Zaïane, and Xindong Wu, editors, *ICDM*, pages 497–506. IEEE, 2011. ISBN 978-0-7695-4408-3.

[28] Arkadiusz Paterek. Improving regularized singular value decomposition for collaborative filtering. In *Proceedings of KDD cup and workshop*, volume 2007, pages 5–8, 2007.

[29] Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme. Bpr: Bayesian personalized ranking from implicit feedback. In *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence*, UAI '09, pages 452–461, Arlington, Virginia, United States, 2009. AUAI Press. ISBN 978-0-9749039-5-8. URL http://dl.acm.org/citation.cfm?id=1795114.1795167.

[30] Ruslan Salakhutdinov and Andriy Mnih. Bayesian probabilistic matrix factorization using markov chain monte carlo. In *Proceedings of the 25th international conference on Machine learning*, pages 880–887. ACM, 2008.

[31] Ruslan Salakhutdinov and Andriy Mnih. Probabilistic matrix factorization. *Advances in neural information processing systems*, 20:1257–1264, 2008.

[32] Badrul M. Sarwar, George Karypis, Joseph A. Konstan, and John T. Riedl. Application of dimensionality reduction in recommender system – a case study. In *IN ACM WEBKDD WORKSHOP*, 2000.

[33] D Seung and L Lee. Algorithms for non-negative matrix factorization. *Advances in neural information processing systems*, 13:556–562, 2001.

[34] Xiaoyuan Su and Taghi M Khoshgoftaar. A survey of collaborative filtering techniques. *Advances in Artificial Intelligence*, 2009:4, 2009.

[35] Mingxuan Sun, Guy Lebanon, and Paul Kidwell. Estimating probabilities in recommendation systems. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 61(3):471–492, 2012.

[36] Kai Yu, Shenghuo Zhu, John Lafferty, and Yihong Gong. Fast nonparametric matrix factorization for large-scale collaborative filtering. In *Proceedings of the 32nd international ACM SIGIR conference on Research and development in information retrieval*, pages 211–218. ACM, 2009.